


# Socket Programming



# Client-Server communication

## ■ Server

- ❑ passively waits for and responds to clients
- ❑ **passive** socket

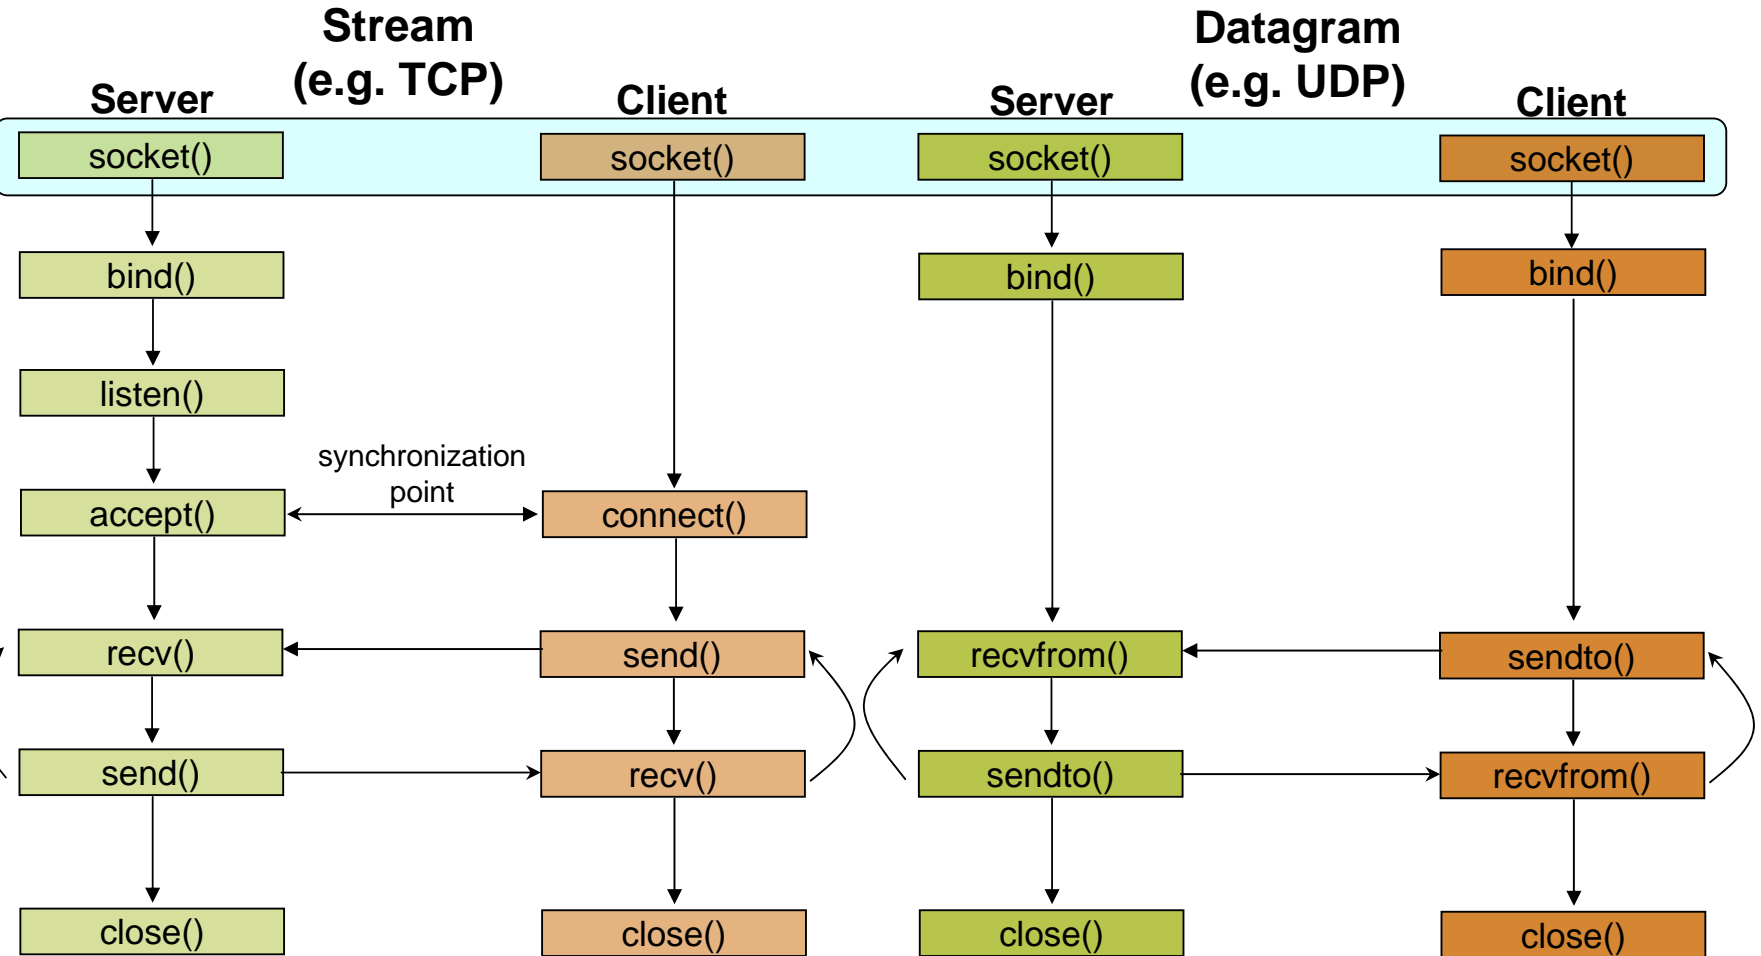
## ■ Client

- ❑ initiates the communication
- ❑ must know the address and the port of the server
- ❑ **active** socket

# Sockets - Procedures

Primitive	Meaning
Socket	Create a new communication endpoint
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection

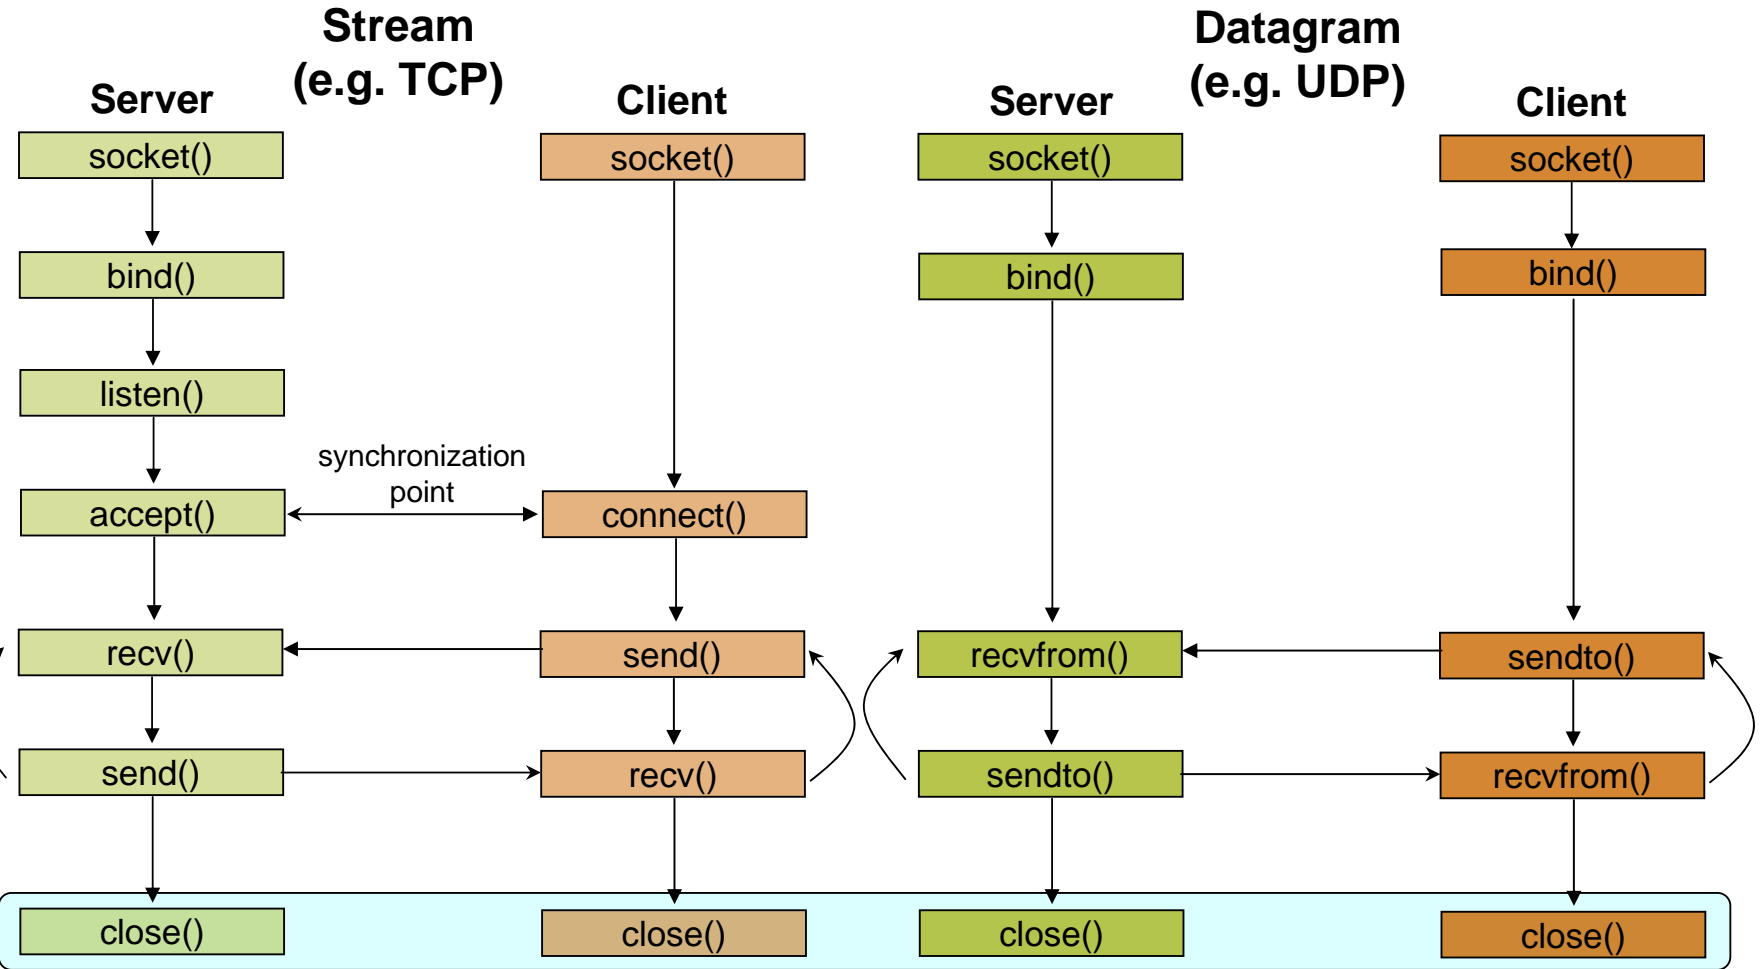
# Client - Server Communication - Unix



# Socket creation in C: `socket ( )`

- `int sockid = socket(family, type, protocol);`
  - ❑ `sockid`: socket descriptor, an integer (like a file-handle)
  - ❑ `family`: integer, communication domain, e.g.,
    - `PF_INET`, IPv4 protocols, Internet addresses (typically used)
    - `PF_UNIX`, Local communication, File addresses
  - ❑ `type`: communication type
    - `SOCK_STREAM` - reliable, 2-way, connection-based service
    - `SOCK_DGRAM` - unreliable, connectionless, messages of maximum length
  - ❑ `protocol`: specifies protocol
    - `IPPROTO_TCP` `IPPROTO_UDP`
    - usually set to 0 (i.e., use default protocol)
  - ❑ upon failure returns -1
- 👉 NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

# Client - Server Communication - Unix



# Socket close in C: `close()`

- When finished using a socket, the socket should be closed
- `status = close(sockid);`
  - ❑ `sockid`: the file descriptor (socket being closed)
  - ❑ `status`: 0 if successful, -1 if error
- Closing a socket
  - ❑ closes a connection (for stream socket)
  - ❑ frees up the port used by the socket

# Specifying Addresses

- Socket API defines a **generic** data type for addresses:

```
struct sockaddr {  
    unsigned short sa_family; /* Address family (e.g. AF_INET) */  
    char sa_data[14];        /* Family-specific address information */  
}
```

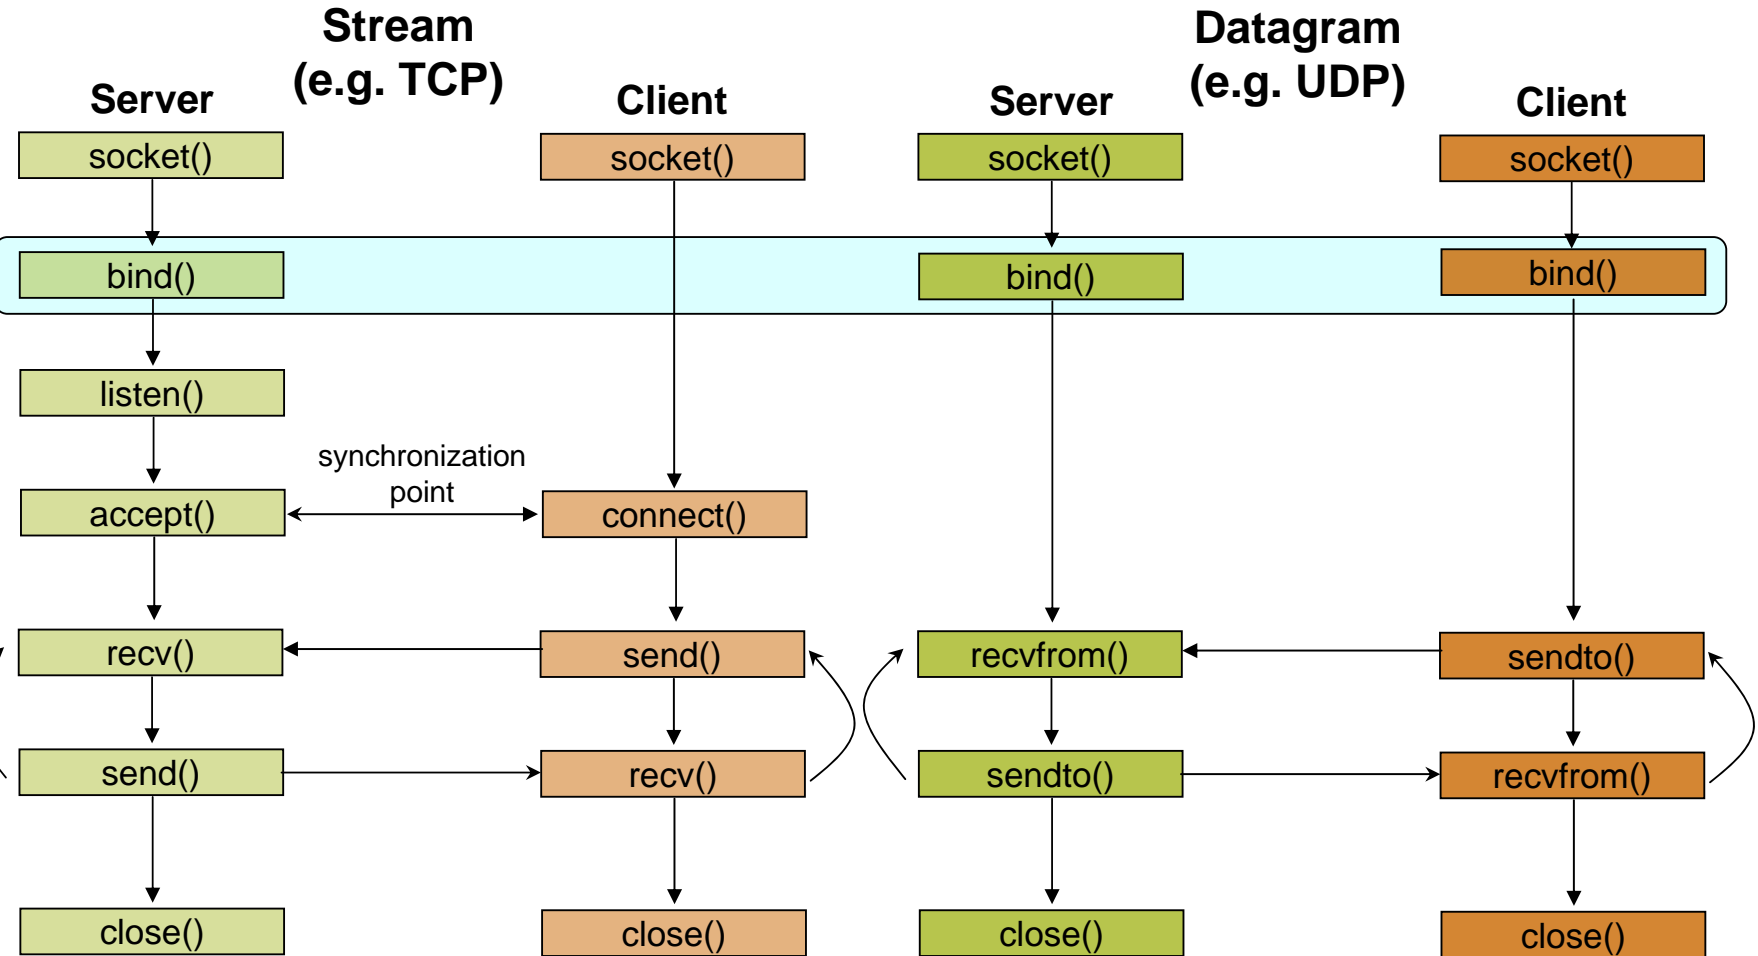
- Particular form of the sockaddr used for **TCP/IP** addresses:

```
struct in_addr {  
    unsigned long s_addr; /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family; /* Internet protocol (AF_INET) */  
    unsigned short sin_port;   /* Address port (16 bits) */  
    struct in_addr sin_addr;   /* Internet address (32 bits) */  
    char sin_zero[8];         /* Not used */  
}
```

👉 **Important:** sockaddr\_in can be casted to a sockaddr



# Client - Server Communication - Unix



# Assign address to socket: `bind()`

- associates and reserves a port for use by the socket
- `int status = bind(sockid, &addrport, size);`
  - `sockid`: integer, socket descriptor
  - `addrport`: struct `sockaddr`, the (IP) address and port of the machine
    - for TCP/IP server, internet address is usually set to `INADDR_ANY`, i.e., chooses any incoming interface
  - `size`: the size (in bytes) of the `addrport` structure
  - `status`: upon failure -1 is returned

# bind( ) - Example with TCP

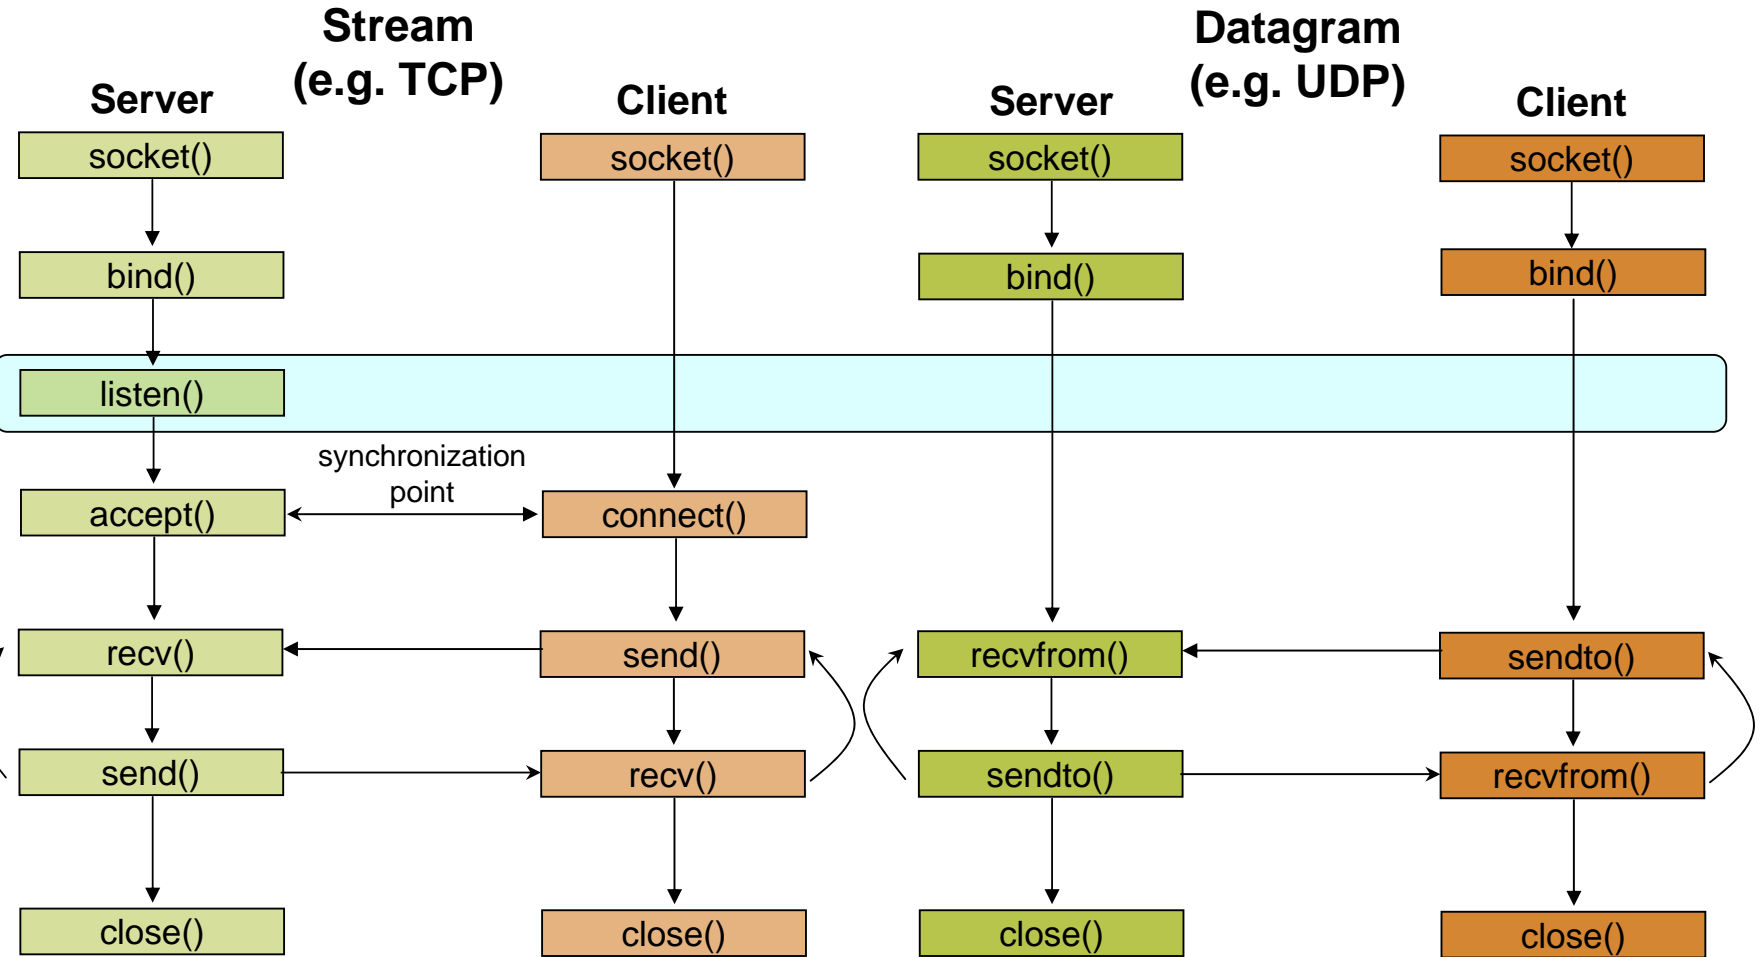
```
int sockid;
struct sockaddr_in addrport;
sockid = socket(PF_INET, SOCK_STREAM, 0);

addrport.sin_family = AF_INET;
addrport.sin_port = htons(5100);
addrport.sin_addr.s_addr = htonl(INADDR_ANY);
if(bind(sockid, (struct sockaddr *) &addrport, sizeof(addrport))!= -1) {
    ...}
```

# Skipping the `bind()`

- `bind` can be skipped for both types of sockets
- Datagram socket:
  - ❑ if only sending, no need to bind. The OS finds a port each time the socket sends a packet
  - ❑ if receiving, need to bind
- Stream socket:
  - ❑ destination determined during connection setup
  - ❑ don't need to know port sending from (during connection setup, receiving end is informed of port)

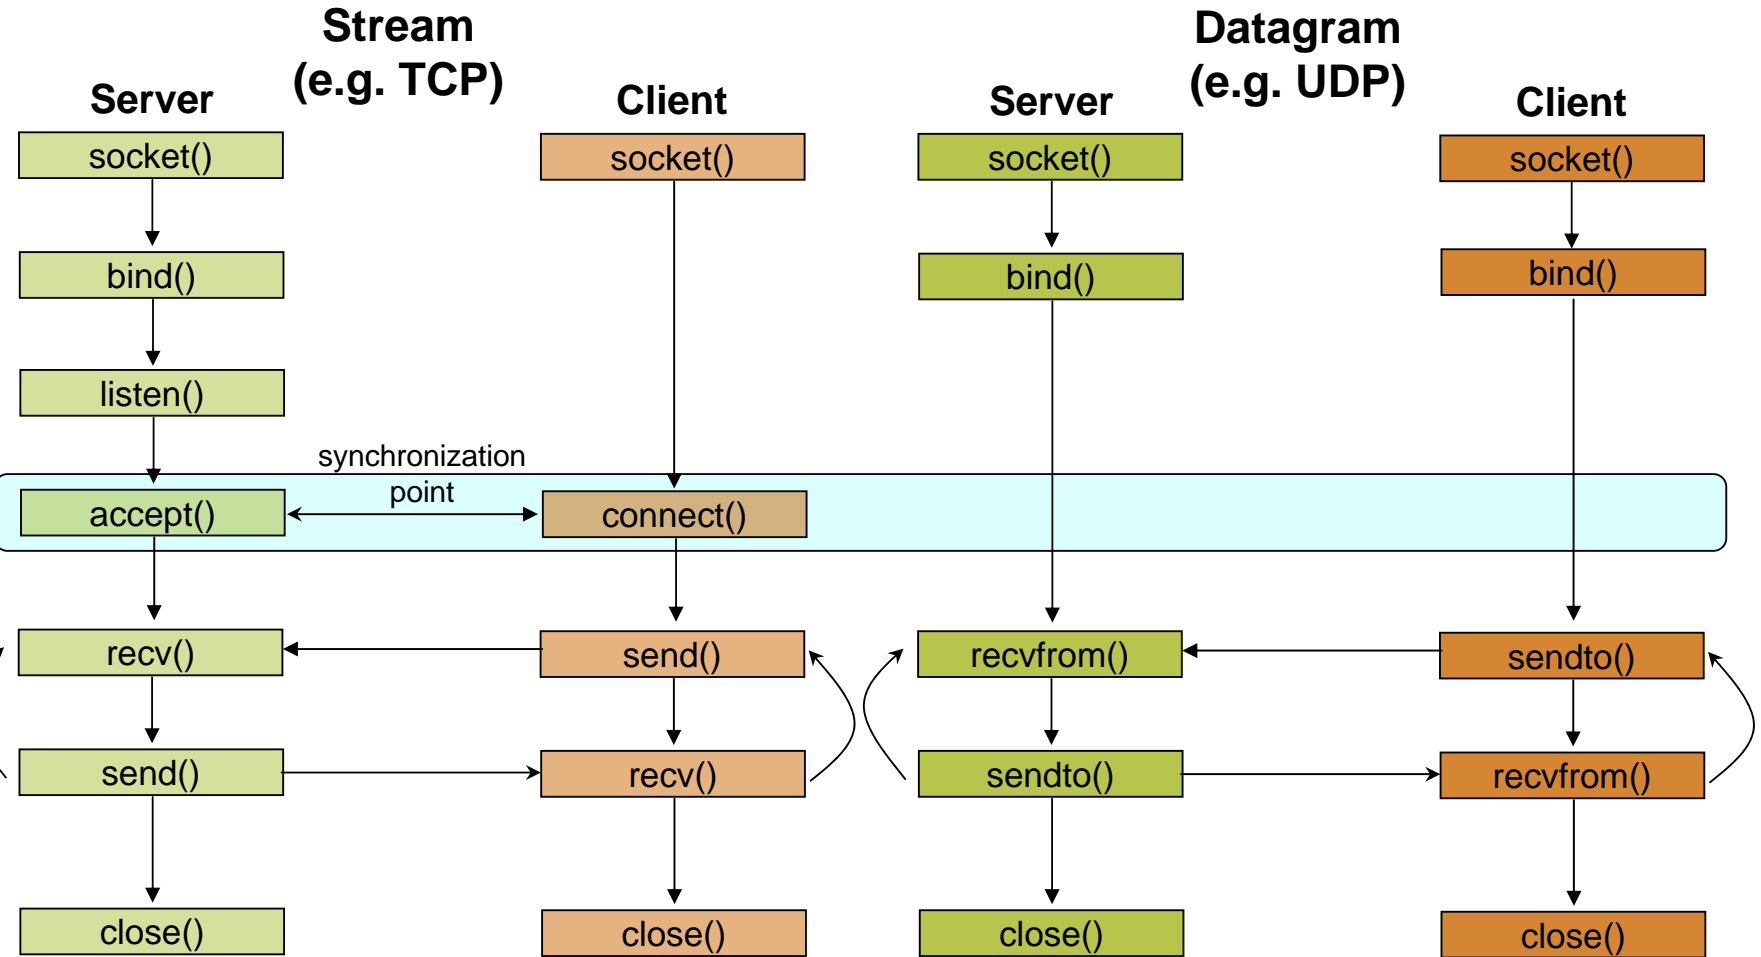
# Client - Server Communication - Unix



# Assign address to socket: `bind( )`

- Instructs TCP protocol implementation to listen for connections
- `int status = listen(sockid, queueLimit);`
  - `sockid`: integer, socket descriptor
  - `queueLen`: integer, # of active participants that can “wait” for a connection
  - `status`: 0 if listening, -1 if error
- `listen( )` is **non-blocking**: returns immediately
- The listening socket (`sockid`)
  - is never used for sending and receiving
  - is used by the server only as a way to get new sockets

# Client - Server Communication - Unix



# Establish Connection: connect ( )

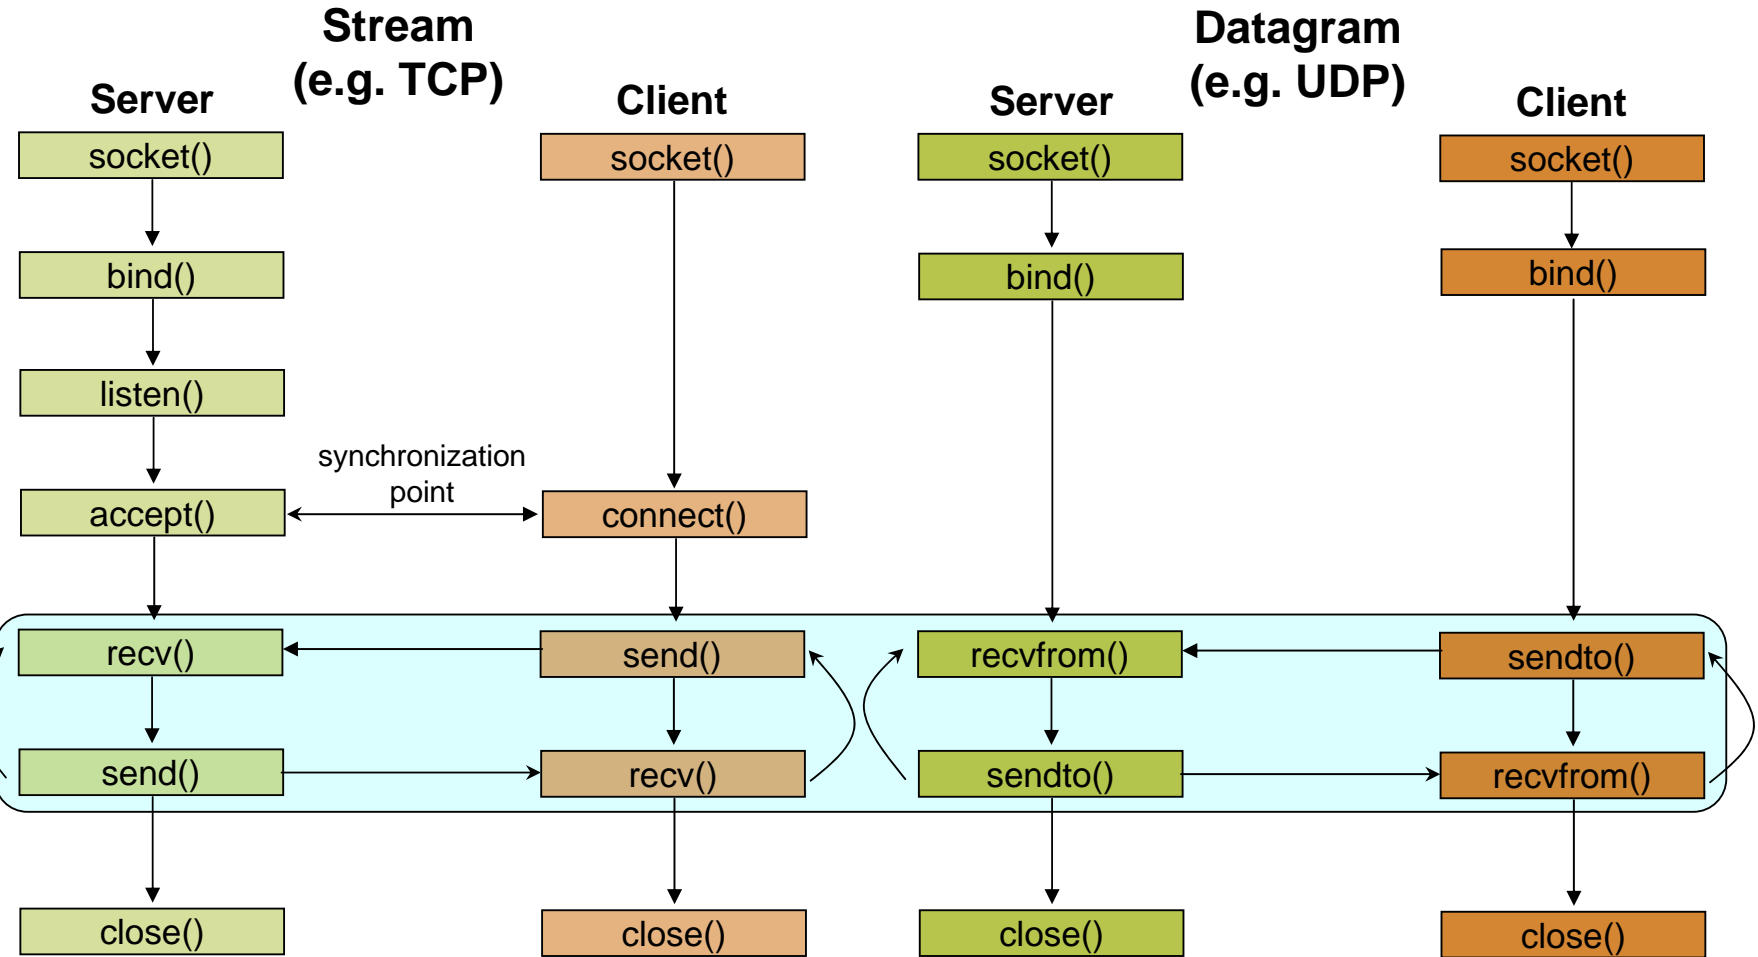
- The client establishes a connection with the server by calling `connect ( )`
- `int status = connect(sockid, &foreignAddr, addrlen);`
  - `sockid`: integer, socket to be used in connection
  - `foreignAddr`: struct `sockaddr`: address of the passive participant
  - `addrlen`: integer, `sizeof(name)`
  - `status`: 0 if successful connect, -1 otherwise
- `connect ( )` is **blocking**



# Incoming Connection: `accept ( )`

- The server gets a socket for an incoming client connection by calling `accept ( )`
- `int s = accept(sockid, &clientAddr, &addrLen);`
  - `s`: integer, the new socket (used for data-transfer)
  - `sockid`: integer, the orig. socket (being listened on)
  - `clientAddr`: struct `sockaddr`, address of the active participant
    - filled in upon return
  - `addrLen`: `sizeof(clientAddr)`: value/result parameter
    - must be set appropriately before call
    - adjusted upon return
- `accept ( )`
  - is **blocking**: waits for connection before returning
  - dequeues the next connection on the queue for socket (`sockid`)

# Client - Server Communication - Unix



# Exchanging data with stream socket

- `int count = send(sockid, msg, msgLen, flags);`
  - ❑ `msg`: `const void[]`, message to be transmitted
  - ❑ `msgLen`: integer, length of message (in bytes) to transmit
  - ❑ `flags`: integer, special options, usually just 0
  - ❑ `count`: # bytes transmitted (-1 if error)
- `int count = recv(sockid, recvBuf, bufLen, flags);`
  - ❑ `recvBuf`: `void[]`, stores received bytes
  - ❑ `bufLen`: # bytes received
  - ❑ `flags`: integer, special options, usually just 0
  - ❑ `count`: # bytes received (-1 if error)
- Calls are **blocking**
  - ❑ returns only after data is sent / received

# Exchanging data with datagram socket

- `int count = sendto(sockid, msg, msgLen, flags, &foreignAddr, addrlen);`
  - `msg, msgLen, flags`, count: same with `send( )`
  - `foreignAddr`: struct `sockaddr`, address of the destination
  - `addrlen`: `sizeof(foreignAddr)`
- `int count = recvfrom(sockid, recvBuf, bufLen, flags, &clientAddr, addrlen);`
  - `recvBuf, bufLen, flags`, count: same with `recv( )`
  - `clientAddr`: struct `sockaddr`, address of the client
  - `addrlen`: `sizeof(clientAddr)`
- Calls are **blocking**
  - returns only after data is sent / received

# Example - Echo

- A client communicates with an “echo” server
- The server simply echoes whatever it receives back to the client